

Bike Share Demand Forecast

November 20, 2019

```
[ ]: # ----- Step 1a: Define and categorize the problem statement
    ↪ -----

# The problem statement is to "Predict the daily bike rental count based on the
    ↪ environmental and seasonal settings"
# This is clearly a 'Supervised machine learning regression problem' to predict
    ↪ a number based on the input features.

# ----- Step 1a ends here -----

[1]: # ----- Step 1b: Import all the required libraries -----

#---- for data transformations----
#install.packages("lubridate")
library(lubridate)

#---- for EDA Visualizations -----
#install.packages("corrplot")
library(corrplot)
#install.packages("ggplot2")
library(ggplot2)
#install.packages("GGally")
library("GGally")
#install.packages("ggExtra")
library(ggExtra)

#---- for model building----
library(caret)
#install.packages("Metrics")
library(Metrics)
#install.packages("randomForest")
library(randomForest)

#install.packages(gbm)
library (gbm)

# ----- Step 1b ends here -----
```

Attaching package: lubridate

The following object is masked from package:base:

date

corrplot 0.84 loaded

Registered S3 method overwritten by 'GGally':

method from

+.gg ggplot2

Loading required package: lattice

Attaching package: Metrics

The following objects are masked from package:caret:

precision, recall

randomForest 4.6-14

Type rfNews() to see new features/changes/bug fixes.

Attaching package: randomForest

The following object is masked from package:ggplot2:

margin

```
[2]: # ----- Step 2: Gather the data -----  
  
# Data is provided as .csv file and already split into Test and Train.  
# The training set is comprised of the first 19 days of each month, while the  
→test set is the 20th to the end of the month.  
# Let's import the data  
bike= read.csv("/Users/snehashrungarpawar/Documents/Master in Data Science/  
→DPA/Project/Data/train.csv", header=TRUE)  
bike_test = read.csv("/Users/snehashrungarpawar/Documents/Master in Data_  
→Science/DPA/Project/Data/test.csv", header=TRUE)  
# ----- Step 2 ends here -----  
  
[: # ----- Step 3: Data Preparation -----  
  
# 3a. Analyze Attributes: Check properties of data  
# 3b. Complete Data Perform missing value analysis and Impute if needed  
# 3c. Correct Data: Check for any invalid data points  
# 3d. Create Derived Attributes - Feature Extraction  
# 3e. Convert - Converting data to proper formats
```

```
[3]: # 3a. Analyze Attributes: Check properties of data
      dim(bike)
      str(bike)
      head(bike, 10)
      # 3a -> Inference:
          #i. The dataset has 10,886 observations (n=10886) and 12 columns of
      → type int, num and factor.
          #ii. Season, Holiday, Working day and weather are categorical variables.
          #ii. temp, atemp, humidity, windspeed, casual, registered and count are
      → continuous numerical variables.
```

1. 10886 2. 12

```
'data.frame': 10886 obs. of 12 variables:
 $ datetime : Factor w/ 10886 levels "2011-01-01 00:00:00",...: 1 2 3 4 5 6 7 8
9 10 ...
 $ season    : int  1 1 1 1 1 1 1 1 1 1 ...
 $ holiday   : int  0 0 0 0 0 0 0 0 0 0 ...
 $ workingday: int  0 0 0 0 0 0 0 0 0 0 ...
 $ weather   : int  1 1 1 1 1 2 1 1 1 1 ...
 $ temp      : num  9.84 9.02 9.02 9.84 9.84 ...
 $ atemp     : num  14.4 13.6 13.6 14.4 14.4 ...
 $ humidity  : int  81 80 80 75 75 75 80 86 75 76 ...
 $ windspeed : num  0 0 0 0 0 ...
 $ casual    : int  3 8 5 3 0 0 2 1 1 8 ...
 $ registered: int  13 32 27 10 1 1 0 2 7 6 ...
 $ count     : int  16 40 32 13 1 1 2 3 8 14 ...
```

	datetime	season	holiday	workingday	weather	temp	atemp	humidity	windspeed
	2011-01-01 00:00:00	1	0	0	1	9.84	14.395	81	0.0000
	2011-01-01 01:00:00	1	0	0	1	9.02	13.635	80	0.0000
	2011-01-01 02:00:00	1	0	0	1	9.02	13.635	80	0.0000
	2011-01-01 03:00:00	1	0	0	1	9.84	14.395	75	0.0000
	2011-01-01 04:00:00	1	0	0	1	9.84	14.395	75	0.0000
	2011-01-01 05:00:00	1	0	0	2	9.84	12.880	75	6.0032
	2011-01-01 06:00:00	1	0	0	1	9.02	13.635	80	0.0000
	2011-01-01 07:00:00	1	0	0	1	8.20	12.880	86	0.0000
	2011-01-01 08:00:00	1	0	0	1	9.84	14.395	75	0.0000
	2011-01-01 09:00:00	1	0	0	1	13.12	17.425	76	0.0000

```
[4]: # 3b. Complete Data Perform missing value analysis and Impute if needed
      table(is.na(bike))
      # 3b -> Inference: There are no null values in the dataset. If it had, then
      → either the rows/columns had to be
          # dropped or the null values be imputed based on the % of null values
```

FALSE
130632

```
[ ]: # 3c. Correct Data: Check for any invalid data points
      # From above observations data doesnot seem to have any invalid datatypes
      →to be handled.
      # Let's check for the outliers in EDA step

[5]: # 3d. Create Derived Attributes - Feature Extraction
      # Lets extract 'date', 'month', 'weekday' and 'year' from 'datetime' column
      →as we will be needing it for analysis

      bike$date=as.factor(day(bike$datetime))
      bike$year = as.factor(year(bike$datetime))
      bike$month = as.factor(month(bike$datetime))
      bike$hour = as.factor(hour(bike$datetime))
      bike$wkday = as.factor(wday(bike$datetime))

      bike_test$date=as.factor(day(bike_test$datetime))
      bike_test$year = as.factor(year(bike_test$datetime))
      bike_test$month = as.factor(month(bike_test$datetime))
      bike_test$hour = as.factor(hour(bike_test$datetime))
      bike_test$wkday = as.factor(wday(bike_test$datetime))

      # Drop datetime as we have extracted all the above needed information
      →from it
      bike = bike[-c(1)]
      bike_test = bike_test[-c(1)]

      head(bike, 5)
      head(bike_test, 5)

      # 3d -> Inference: There are no null values in the dataset. If it had, then
      →either the rows/columns had to be
                        #dropped or the null values be imputed based on the % of
      →null values.
```

season	holiday	workingday	weather	temp	atemp	humidity	windspeed	casual	registered		
1	0	0	1	9.84	14.395	81	0	3	13		
1	0	0	1	9.02	13.635	80	0	8	32		
1	0	0	1	9.02	13.635	80	0	5	27		
1	0	0	1	9.84	14.395	75	0	3	10		
1	0	0	1	9.84	14.395	75	0	0	1		
season	holiday	workingday	weather	temp	atemp	humidity	windspeed	date	year	month	
1	0	1	1	10.66	11.365	56	26.0027	20	2011	1	
1	0	1	1	10.66	13.635	56	0.0000	20	2011	1	
1	0	1	1	10.66	13.635	56	0.0000	20	2011	1	
1	0	1	1	10.66	12.880	56	11.0014	20	2011	1	
1	0	1	1	10.66	12.880	56	11.0014	20	2011	1	

```
[6]: # 3e. Convert - Converting data to proper formats
```

```

# We can clearly see that "season", "holiday", "workingday" and "weather"
→ are categories rather than continuous variable.

# Let's convert them to categories
names = c("season", "holiday", "workingday", "weather")
bike[,names] = lapply(bike[,names], factor)
bike_test[,names] = lapply(bike_test[,names], factor)

str(bike)
str(bike_test)

# ----- Step 3: Data Preparation ends here
→ -----

```

```


'data.frame': 10886 obs. of 16 variables:
 $ season      : Factor w/ 4 levels "1","2","3","4": 1 1 1 1 1 1 1 1 1 1 ...
 $ holiday     : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
 $ workingday  : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
 $ weather     : Factor w/ 4 levels "1","2","3","4": 1 1 1 1 1 2 1 1 1 1 ...
 $ temp       : num  9.84 9.02 9.02 9.84 9.84 ...
 $ atemp      : num  14.4 13.6 13.6 14.4 14.4 ...
 $ humidity   : int   81 80 80 75 75 75 80 86 75 76 ...
 $ windspeed  : num   0 0 0 0 0 ...
 $ casual     : int   3 8 5 3 0 0 2 1 1 8 ...
 $ registered : int  13 32 27 10 1 1 0 2 7 6 ...
 $ count      : int  16 40 32 13 1 1 2 3 8 14 ...
 $ date       : Factor w/ 19 levels "1","2","3","4",...: 1 1 1 1 1 1 1 1 1 1 ...
 $ year       : Factor w/ 2 levels "2011","2012": 1 1 1 1 1 1 1 1 1 1 ...
 $ month      : Factor w/ 12 levels "1","2","3","4",...: 1 1 1 1 1 1 1 1 1 1 ...
 $ hour       : Factor w/ 24 levels "0","1","2","3",...: 1 2 3 4 5 6 7 8 9 10 ...
 $ wkday      : Factor w/ 7 levels "1","2","3","4",...: 7 7 7 7 7 7 7 7 7 7 ...
'data.frame': 6493 obs. of 13 variables:
 $ season      : Factor w/ 4 levels "1","2","3","4": 1 1 1 1 1 1 1 1 1 1 ...
 $ holiday     : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
 $ workingday  : Factor w/ 2 levels "0","1": 2 2 2 2 2 2 2 2 2 2 ...
 $ weather     : Factor w/ 4 levels "1","2","3","4": 1 1 1 1 1 1 1 1 1 2 ...
 $ temp       : num  10.7 10.7 10.7 10.7 10.7 ...
 $ atemp      : num  11.4 13.6 13.6 12.9 12.9 ...
 $ humidity   : int   56 56 56 56 56 60 60 55 55 52 ...
 $ windspeed  : num  26 0 0 11 11 ...
 $ date       : Factor w/ 12 levels "20","21","22",...: 1 1 1 1 1 1 1 1 1 1 ...
 $ year       : Factor w/ 2 levels "2011","2012": 1 1 1 1 1 1 1 1 1 1 ...
 $ month      : Factor w/ 12 levels "1","2","3","4",...: 1 1 1 1 1 1 1 1 1 1 ...
 $ hour       : Factor w/ 24 levels "0","1","2","3",...: 1 2 3 4 5 6 7 8 9 10 ...
 $ wkday      : Factor w/ 7 levels "1","2","3","4",...: 5 5 5 5 5 5 5 5 5 5 ...

```


```
[ ]:
```

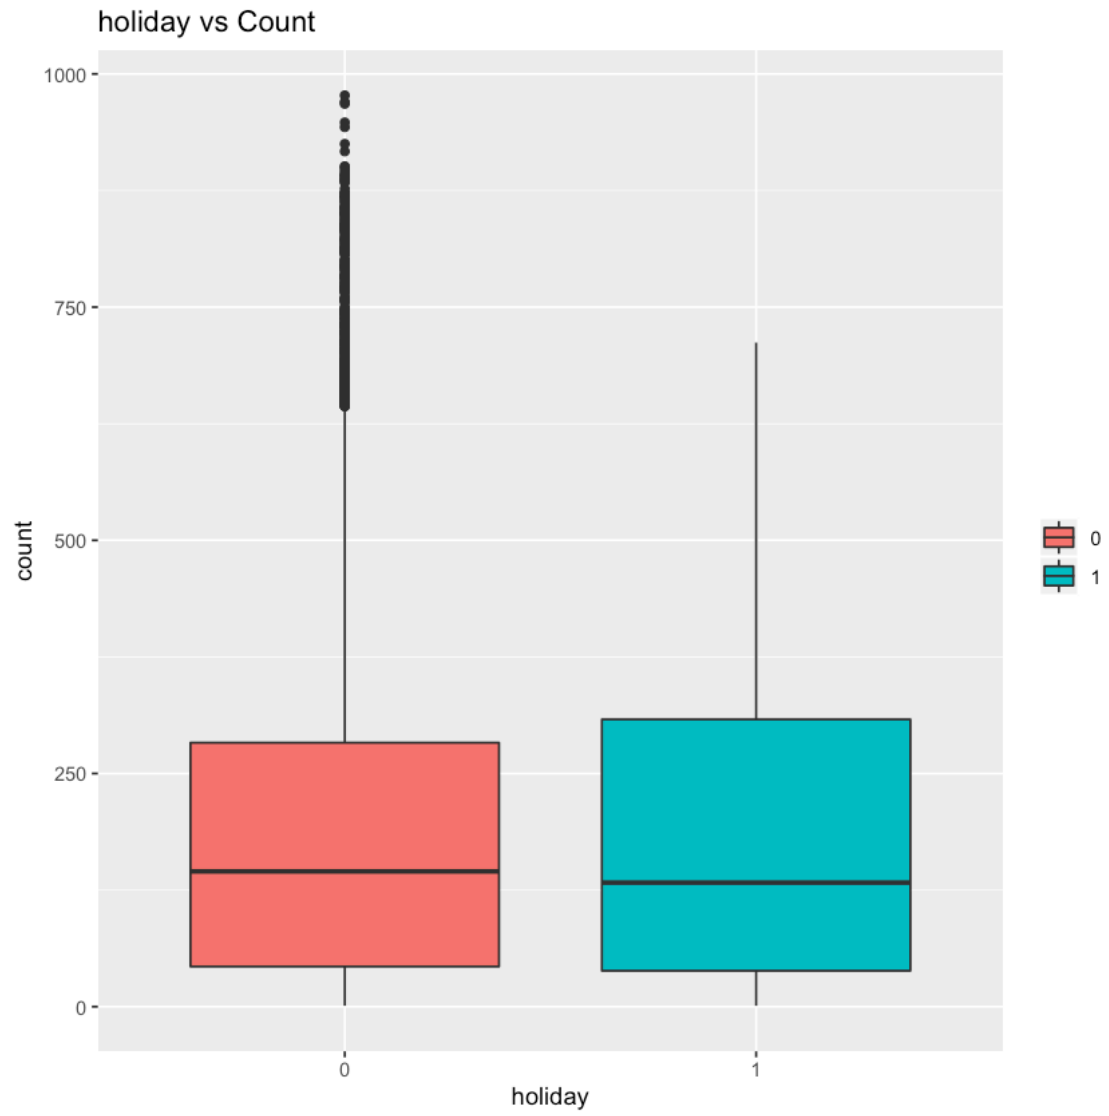
```
[57]: # ----- Step 4: Exploratory Data Analysis -----
      # 4a. Outlier Analysis

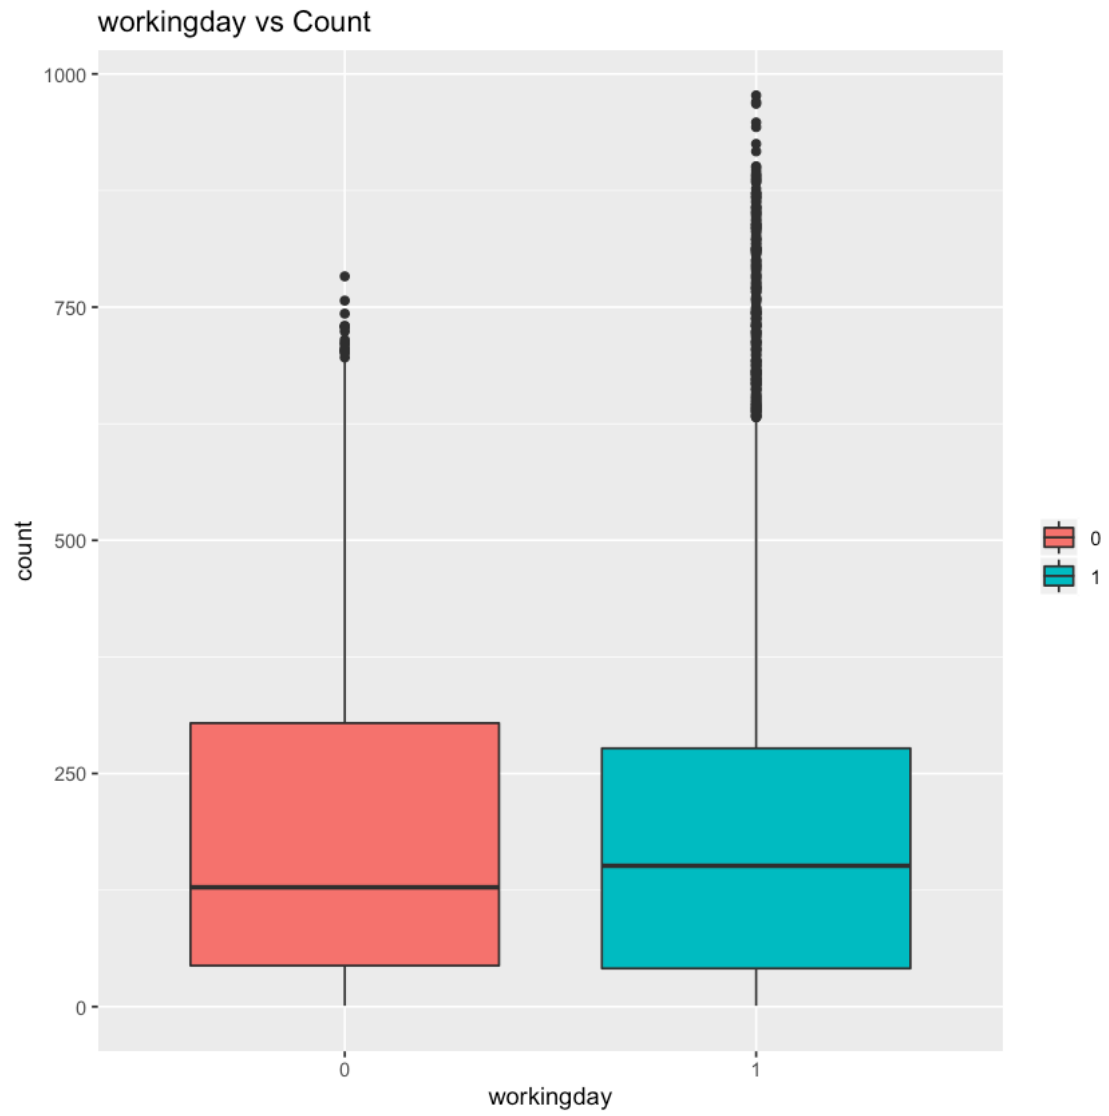
      # 4a(1). Visualize continuos variables
      par(mfrow=c(1,5))
      boxplot(bike$count, main="Count", col="Gray", border = "black")
      boxplot(bike$temp, main="Temperature", col="blue", border = "black")
      boxplot(bike$atemp, main="Feels Like Temp", col="purple", border = "black")
      boxplot(bike$humidity, main="Humidity", col="green", border = "black")
      boxplot(bike$windspeed, main="Windspeed", col="orange", border = "black")
```




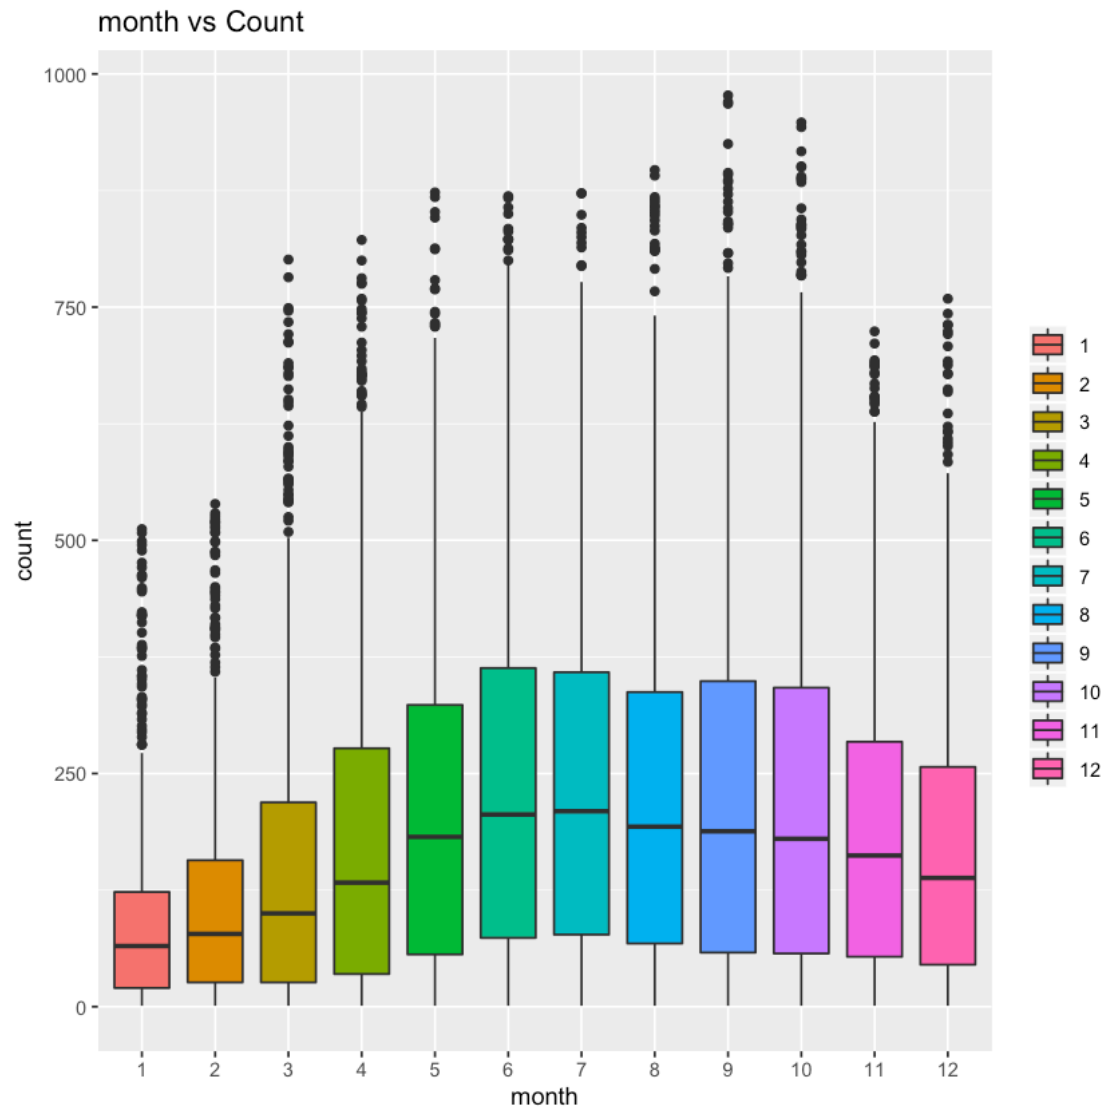
```
[60]: # 4a(2). Visualize categorical variables wrt target variable
      par(mfrow=c(3,4))
      ggplot(data = bike, aes(x=season, y=count, fill=as.factor(season))) +
      ↪geom_boxplot() + labs(title="Season vs Count") + theme(legend.title =
      ↪element_blank())
      ggplot(data = bike, aes(x=weather, y=count, fill=as.factor(weather))) +
      ↪geom_boxplot() + labs(title="weather vs Count") + theme(legend.title =
      ↪element_blank())
      ggplot(data = bike, aes(x=holiday, y=count, fill=as.factor(holiday))) +
      ↪geom_boxplot() + labs(title="holiday vs Count") + theme(legend.title =
      ↪element_blank())
      ggplot(data = bike, aes(x=workingday, y=count, fill=as.factor(workingday)))
      ↪+ geom_boxplot() + labs(title="workingday vs Count") + theme(legend.title =
      ↪element_blank())
      ggplot(data = bike, aes(x=year, y=count, fill=as.factor(year))) +
      ↪geom_boxplot() + labs(title="year vs Count") + theme(legend.title =
      ↪element_blank())
      ggplot(data = bike, aes(x=month, y=count, fill=as.factor(month))) +
      ↪geom_boxplot() + labs(title="month vs Count") + theme(legend.title =
      ↪element_blank())
      ggplot(data = bike, aes(x=wkday, y=count, fill=as.factor(wkday))) +
      ↪geom_boxplot() + labs(title="weekday vs Count") + theme(legend.title =
      ↪element_blank())
      ggplot(data = bike, aes(x=hour, y=count, fill=as.factor(hour))) +
      ↪geom_boxplot() + labs(title="hour vs Count") + theme(legend.title =
      ↪element_blank())
      ggplot(data = bike, aes(x=date, y=count, fill=as.factor(day(date)))) +
      ↪geom_boxplot() + labs(title="date vs Count") + theme(legend.title =
      ↪element_blank())
```



















```
[ ]: # 4b. Correlation Analysis


# ----- Explore Continuous Variables-----
# 4b(1). Explore continous features
# i. Check distribution of target variable
# ii. Explore correlation between independent continuous variables with
→target variable
# iii. Plot heatmap for correlation matrix (to check for
→multicollinearity)
# iv. Visualize the relationship among all continuous variables using
→pairplots
# v. Explore relationship between independent continuous variables and
→dependent variables using Joint Plot
```


```
[61]: # 4b(1) i. Check distribution of target variable
      hist(bike$count, col="blue")
      plot(bike$count)
      # Inference: Target variable "count" is almost normally distributed.
```







```
[62]: # 4b(1) ii. Explore correlation between independent continuous variables with
      ↪ target variable
      plot(bike$temp,bike$count)
      plot(bike$atemp,bike$count)
      plot(bike$windspeed,bike$count)
      plot(bike$humidity,bike$count)
```






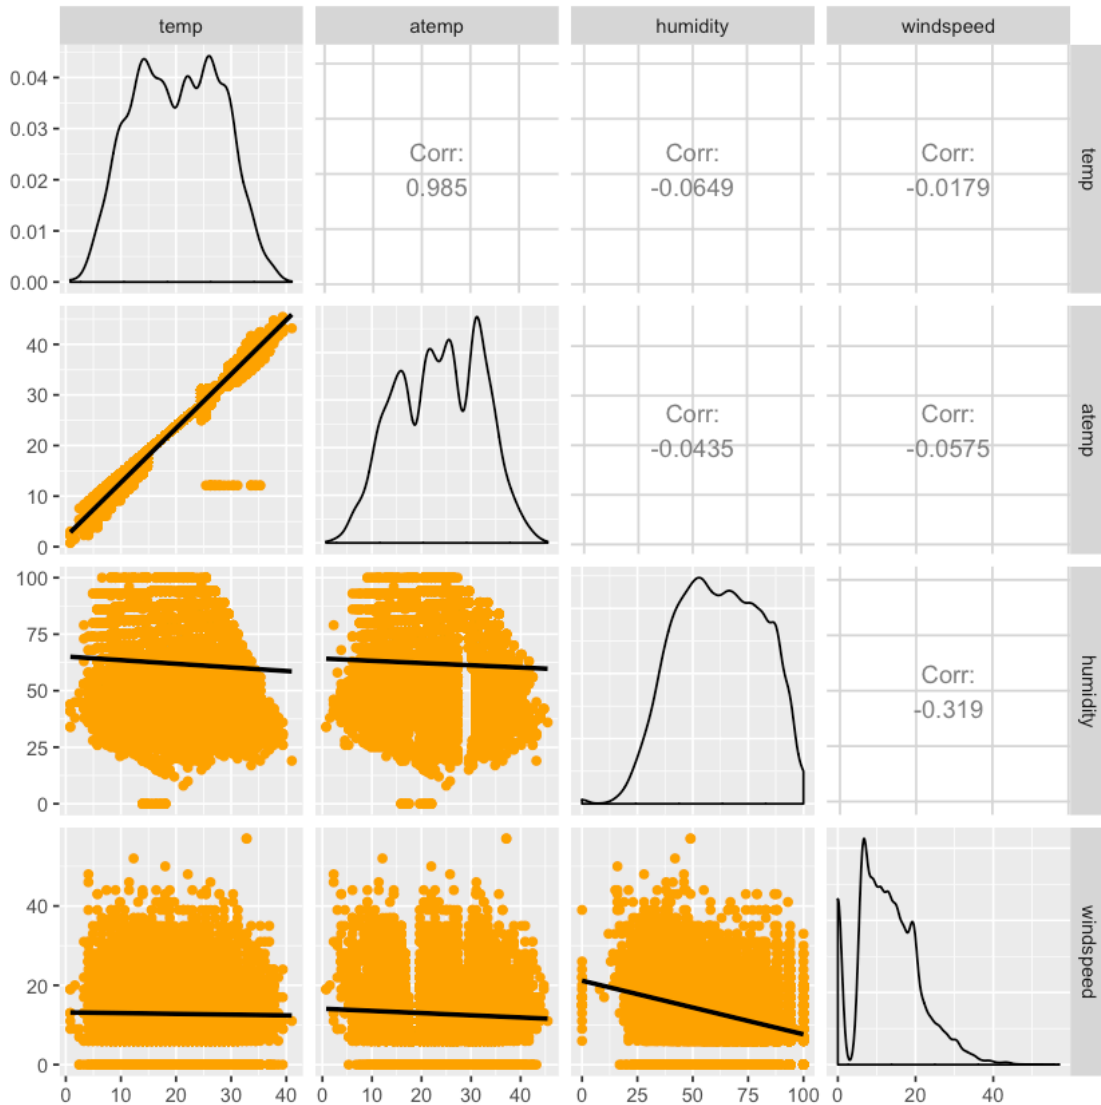




```
[66]: # 4b(1) iii. Plot heatmap for correlation matrix (to check for
→multicollinearity)
      corr <- as.data.frame(lapply(bike[c(5:8, 11)], as.numeric))
      corrplot(cor(corr), method = "color", type='lower')
      # Inference:
      # i. temp and atemp are highly correlated, we would need to drop one of
→them to remove multicollinearity.
      # ii. We can also drop Registered and Casual from our analysis as
→Counts are categorized as Registered and Casual
      # and we will be predicting "Count" variable only.
```




```
[67]: # 4b(1) iv. Visualize the relationship among all continuous variables using
      ↪ pairplots
      ggpairs(bike[c(5:8)], lower=list(continuous=wrap("smooth",
      ↪ colour="orange"))) )
```




[68]: # 4b(1) v. Explore relationship between independent continuous variables and
 → dependent variables using Joint Plot


```
# 1. temp vs Count
plot_center = ggplot(bike, aes(x=temp,y=count)) +
  → geom_point(colour="blue") + geom_smooth(method="lm", colour="grey")
  ggMarginal(plot_center, type="densigram", colour="blue")
# Inference: temp has good correlation with count.
```




```
[69]: # 4b(1).v.2. atemp vs Count
      plot_center = ggplot(bike, aes(x=atemp,y=count)) +
      ↪geom_point(colour="red") + geom_smooth(method="lm", colour="grey")
      ggMarginal(plot_center, type="densigram", colour="red")
      # Inference: atemp has good correlation with count.
```

```
[70]: # 4b(1).v.3. humidity vs Count
      plot_center = ggplot(bike, aes(x=humidity,y=count)) +
      ↪geom_point(colour="green") + geom_smooth(method="lm") +
      ↪geom_smooth(method="lm", colour="grey")
      ggMarginal(plot_center, type="densigram", colour="green")
      # Inference: Humidity has low correlation with count.
```



```
[71]: # 4b(1).v.4. windspeed vs Count
      plot_center = ggplot(bike, aes(x=windspeed,y=count)) +
      ↪geom_point(colour="orange") + geom_smooth(method="lm", colour="grey")
      ggMarginal(plot_center, type="densigram", colour="orange")
```



```
[ ]: # 4b(1) Inferences Summary - Analysis of continuous variables
      # 1. Target variable 'count' is almost normally distributed.
      # 2. From correlation with dependent variable "count", we can see that
      → 'casual', 'registered' are very
         # highly correlated to cnt. Needs to be dropped from the dataset.
      # 3. 'humidity' has low correlation with 'count'. For now, let's keep it.
      # 4. atemp and temp has good correlation with 'count'
      # 5. From heatmap, we can see that atemp and temp are highly correlated.
      → So we need to drop 1 to remove multicollinearity.
         # 6. Since, as seen from jointplot,  $p(\text{atemp}) < p(\text{temp})$ , we can drop
      → 'temp' and retain 'atemp' in the dataset.
```

```

[72]: # ----- Explore Catogorical Variables-----
      # 4b(2) Explore categorical features
      # i. Check distribution of categorical variables
      ggplot(bike, aes(x=" ",fill=year))+ geom_bar(width = 1)+  

→coord_polar("y")+labs(title = "year")+theme_void()

      ggplot(bike, aes(x=" ",fill=month))+ geom_bar(width = 1)+  

→coord_polar("y")+labs(title = "month")+theme_void()
      bike$season = factor(bike$season)

      ggplot(bike, aes(x=" ",fill=season))+ geom_bar(width = 1)+  

→coord_polar("y")+labs(title = "Season")+theme_void()
      bike$holiday = factor(bike$holiday)

      ggplot(bike, aes(x=" ",fill=holiday))+ geom_bar(width = 1)+  

→coord_polar("y")+labs(title = "holiday")+theme_void()

      ggplot(bike, aes(x=" ",fill=wkday))+ geom_bar(width = 1)+  

→coord_polar("y")+labs(title = "weekday")+theme_void()
      bike$workingday = factor(bike$workingday)

      ggplot(bike, aes(x=" ",fill=workingday))+ geom_bar(width = 1)+  


→coord_polar("y")+labs(title = "workingday")+theme_void()
      bike$weather = factor(bike$weather)

      ggplot(bike, aes(x=" ",fill=weather))+ geom_bar(width = 1)+  


→coord_polar("y")+labs(title = "weather")+theme_void()

```


year




month



Season




holiday




weekday



workingday



weather



```
[73]: # ii. Check how individual categorical features affects the target variable
      ggplot(bike, aes(x=season, y=count, fill=year)) +
        stat_summary(
          fun.y=median,
          geom='bar',
          position=position_dodge(),
        ) + labs(title="Histogram for Seasons") + labs(x="Season", y="Count")

      ggplot(bike, aes(x=year, y=count, fill=year)) +
        stat_summary(
          fun.y=median,
          geom='bar',
```

```

        position=position_dodge(),
    ) +
    labs(title="Histogram for year") + labs(x="year", y="Count")

ggplot(bike, aes(x=month, y=count, fill=month)) +
  stat_summary(
    fun.y=median,
    geom='bar',
    position=position_dodge(),
  ) +
  labs(title="Histogram for month") + labs(x="month", y="Count")

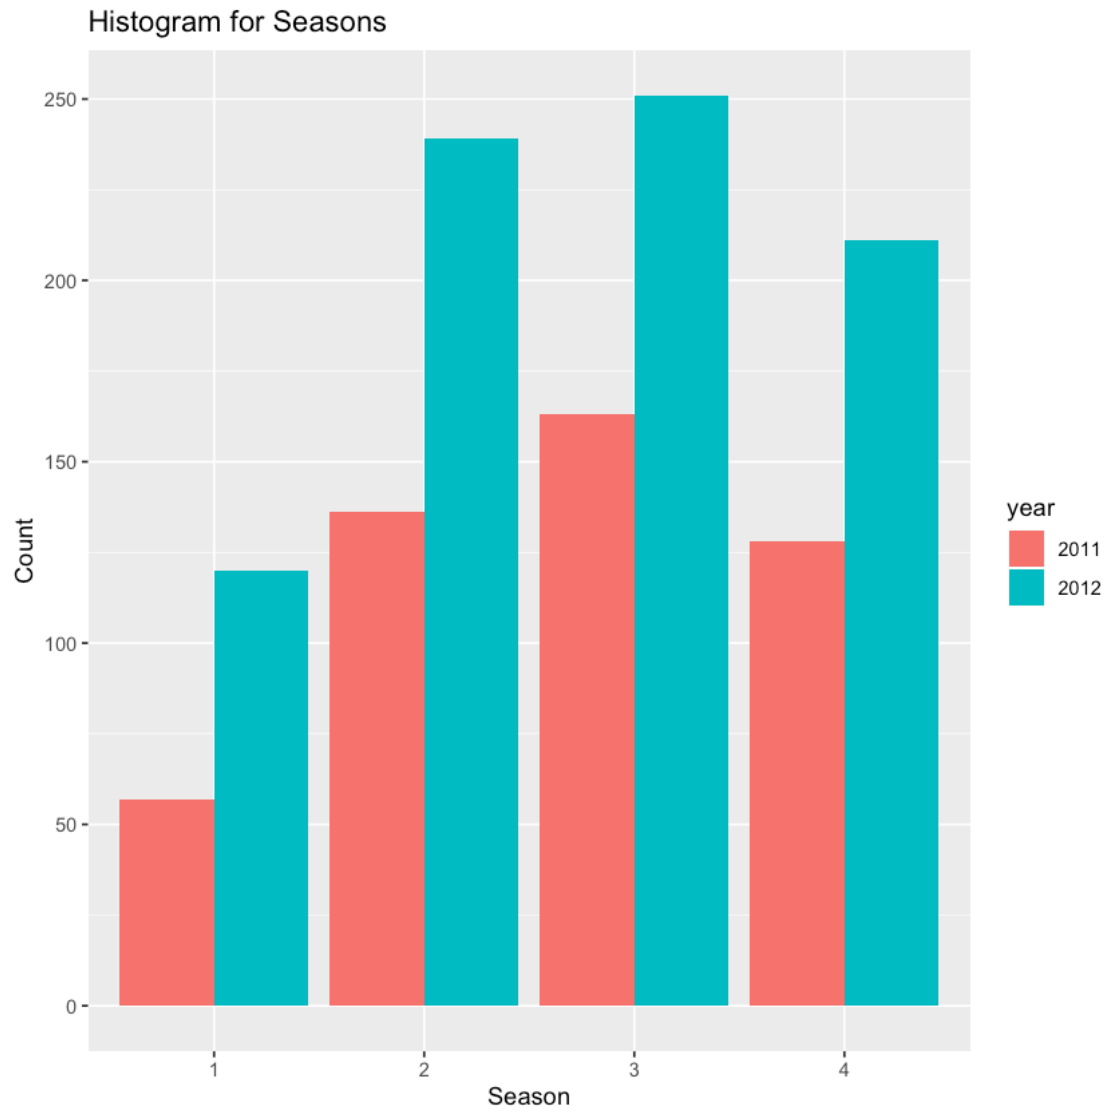
ggplot(bike, aes(x=holiday, y=count, fill=holiday)) +
  stat_summary(
    fun.y=median,
    geom='bar',
    position=position_dodge(),
  ) + labs(title="Histogram for holiday") +labs(x="holiday", y="Count")


ggplot(bike, aes(x=wkday, y=count, fill=wkday)) +
  stat_summary(
    fun.y=median,
    geom='bar',
    position=position_dodge(),
  ) + labs(title="Histogram for weekday") +labs(x="weekday", y="Count")


ggplot(bike, aes(x=workingday, y=count, fill=workingday)) +
  stat_summary(
    fun.y=median,
    geom='bar',
    position=position_dodge(),
  ) + labs(title="Histogram for working day") +labs(x="working day", y="Count")


ggplot(bike, aes(x=weather, y=count, fill=weather)) +
  stat_summary(
    fun.y=median,
    geom='bar',
    position=position_dodge(),
  ) + labs(title="Histogram for weather") +labs(x="weather", y="Count")


```

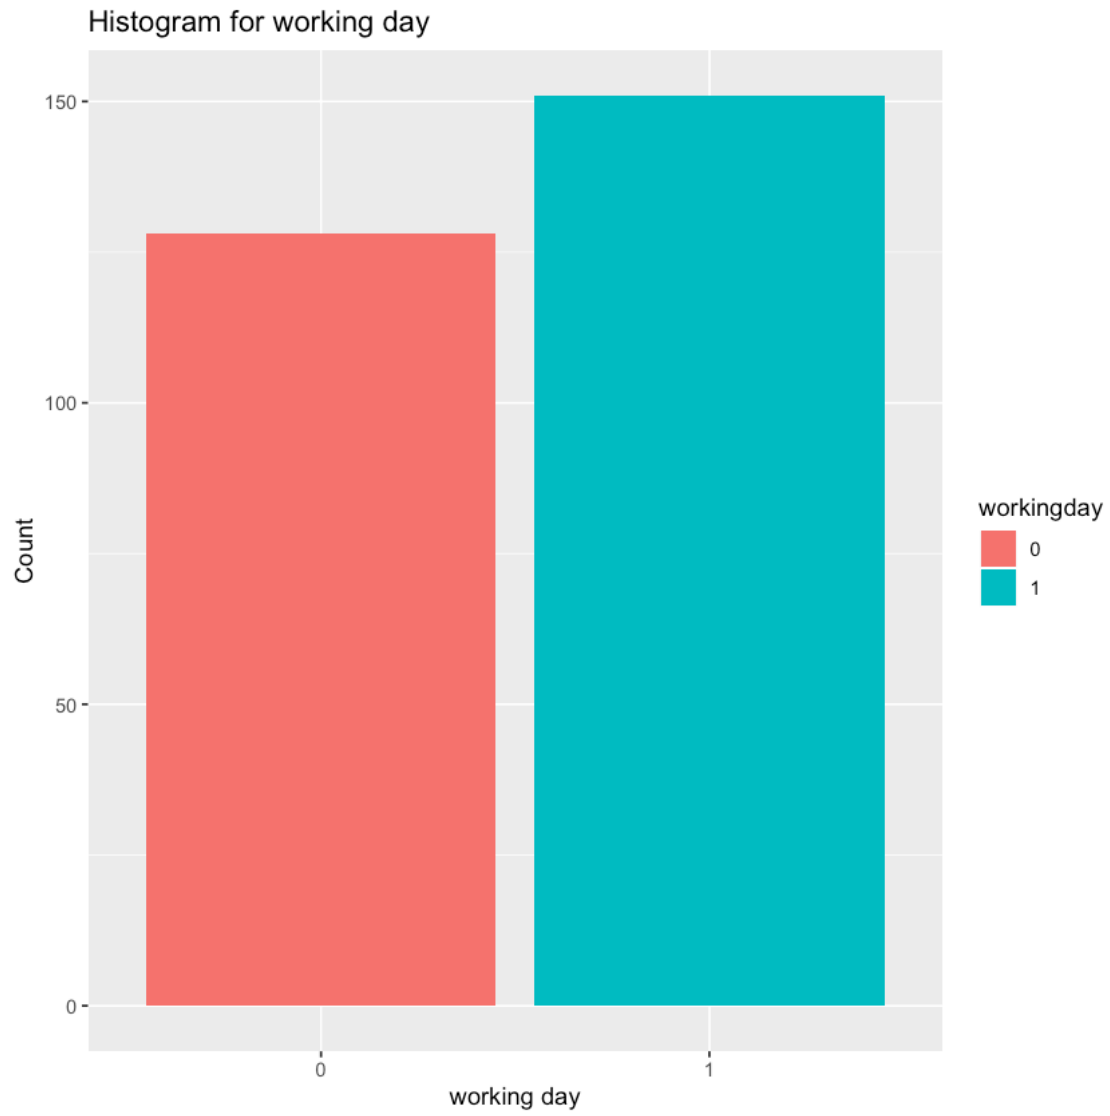















```
[52]: # iii. Explore trends over time ---- exploring some more pairplots

      ggplot(bike, aes(x=season, y=count, group=year, color=year)) +
        stat_summary(
          fun.y=mean,
          geom='line'
        ) +
        stat_summary(
          fun.y=mean,
          geom='point'
        ) +
        labs(title="Average Count by Month Across Season") +
        labs(x="Season", y="Count")
```


```

→+
ggplot(bike, aes(x=bike$hour, y=count, group=season, color=season))
  stat_summary(
    fun.y=mean,
    geom='line'
  ) +
  stat_summary(
    fun.y=mean,
    geom='point'
  )+
  labs(title="Average Count By Hour Of The Day Across Season") +
  labs(x="Hour of the Day", y="Count")


ggplot(bike, aes(x=bike$hour, y=count, group=wkday, color=wkday)) +
  stat_summary(
    fun.y=mean,
    geom='line'
  ) +
  stat_summary(
    fun.y=mean,
    geom='point'
  )+
  labs(title="Average Count By Hour Of The Day Across Weekdays") +
  labs(x="Hour of the Day", y="Count")

ggplot(bike, aes(x=bike$day, y=count, group=day, color=day)) +
  stat_summary(
    fun.y=mean,
    geom='line'
  ) +
  stat_summary(
    fun.y=mean,
    geom='point'
  )+
  labs(title="Average Count By Day") +
  labs(x="HDay", y="Count")


```



Average Count By Hour Of The Day Across Season



Average Count By Hour Of The Day Across Weekdays



```
[74]: # 4c. Drop some variables from the dataset based on the analysis so far
      # drop temp, casual, registered and date
      bike_subset = bike[-c(5,9:10, 12)]
      head(bike_subset,5)
```

season	holiday	workingday	weather	atemp	humidity	windspeed	count	year	month	hour
1	0	0	1	14.395	81	0	16	2011	1	0
1	0	0	1	13.635	80	0	40	2011	1	1
1	0	0	1	13.635	80	0	32	2011	1	2
1	0	0	1	14.395	75	0	13	2011	1	3
1	0	0	1	14.395	75	0	1	2011	1	4

```
[ ]: #----- Step 4: Exploratory Data Analysis ENDS Here-----
      # Final observations:
```

```

#1.) 'atemp' and 'temp' are very strongly correlated . Drop 'atemp' from the
→dataset (since it has higher p-value
    #than 'temp')
#2.) 'date' does not seem to have any affect on count of bikes, it can be
→dropped from the dataset
#-----

```

```

[ ]: #-----Part 5 : Model Builing starts here -----
      # 5a. Split data into test and train set
      # 5b. Linear Regression
      # 5c. Random Forest
      # 5d. Gradient Boosting

```

```

[7]: # 5a. Split data into test and train set
      sample_size = floor(0.8 * nrow(bike))
      set.seed(1)
      train_index = sample(nrow(bike), size = sample_size)
      train <- bike[train_index, ]
      test <- bike[-train_index, ]

```

```

[8]: # 5b. Linear Regression
      # Fit Linear Model
      # drop atemp, registered, casual and date
      train_subset = train[-c(6,9:10, 12)]
      test_subset = test[-c(6,9:10, 12)]

      lm_fit = lm(count ~ ., data = train_subset)
      summary(lm_fit)

      # Choosing the best model by AIC in a Stepwise Algorithm
      # The step() function iteratively removes insignificant features from
→the model.
      step(lm_fit)
      summary(lm_fit)

      # Calculate Train RMSLE
      y_act_train <- abs(train_subset$count)
      y_pred_train <- abs(predict(lm_fit, train_subset))
      lm_train_RMSLE = rmsle(y_act_train, y_pred_train)

      # Calculate Test RMSLE
      y_act_test <- abs(test_subset$count)
      y_pred_test <- abs(predict(lm_fit, test_subset))
      lm_test_RMSLE = rmsle(y_act_test, y_pred_test)

      # Save the results
      lm_results = predict(lm_fit, bike_test)
      hist(lm_results)

```


Call:

```
lm(formula = count ~ ., data = train_subset)
```

Residuals:

Min	1Q	Median	3Q	Max
-351.68	-61.47	-7.12	50.93	438.13

Coefficients: (4 not defined because of singularities)

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-84.59331	9.33850	-9.059	< 2e-16	***
season2	67.39603	7.87769	8.555	< 2e-16	***
season3	76.82918	7.64092	10.055	< 2e-16	***
season4	75.49214	5.62907	13.411	< 2e-16	***
holiday1	11.78025	7.81568	1.507	0.131781	
workingday1	14.71109	4.06430	3.620	0.000297	***
weather2	-12.58747	2.67528	-4.705	2.58e-06	***
weather3	-71.11945	4.47545	-15.891	< 2e-16	***
weather4	-174.84435	100.78681	-1.735	0.082813	.
temp	5.03023	0.33352	15.082	< 2e-16	***
humidity	-0.76332	0.07849	-9.724	< 2e-16	***
windspeed	-0.54238	0.14408	-3.764	0.000168	***
year2012	86.95570	2.19273	39.656	< 2e-16	***
month2	11.13212	5.39034	2.065	0.038934	*
month3	29.41766	5.77185	5.097	3.53e-07	***
month4	-16.76422	5.99977	-2.794	0.005215	**
month5	12.64662	5.45402	2.319	0.020431	*
month6	NA	NA	NA	NA	
month7	-36.80569	5.57965	-6.596	4.46e-11	***
month8	-27.30585	5.45357	-5.007	5.64e-07	***
month9	NA	NA	NA	NA	
month10	21.16010	5.77289	3.665	0.000248	***
month11	1.14471	5.35424	0.214	0.830712	
month12	NA	NA	NA	NA	
hour1	-11.39400	7.48756	-1.522	0.128115	
hour2	-24.02145	7.45155	-3.224	0.001270	**
hour3	-37.44770	7.55542	-4.956	7.32e-07	***
hour4	-38.01239	7.44329	-5.107	3.34e-07	***
hour5	-23.47057	7.47555	-3.140	0.001697	**
hour6	36.59158	7.41431	4.935	8.15e-07	***
hour7	170.52864	7.39633	23.056	< 2e-16	***
hour8	311.38508	7.44159	41.844	< 2e-16	***
hour9	164.73930	7.38202	22.316	< 2e-16	***
hour10	113.53297	7.46630	15.206	< 2e-16	***
hour11	140.80547	7.50670	18.757	< 2e-16	***
hour12	177.90103	7.56740	23.509	< 2e-16	***
hour13	177.19756	7.66452	23.119	< 2e-16	***
hour14	162.02489	7.65093	21.177	< 2e-16	***
hour15	168.32943	7.59391	22.166	< 2e-16	***

hour16	231.47403	7.62564	30.355	< 2e-16	***
hour17	387.32373	7.66555	50.528	< 2e-16	***
hour18	360.43568	7.58264	47.534	< 2e-16	***
hour19	245.60360	7.43058	33.053	< 2e-16	***
hour20	164.11798	7.51229	21.847	< 2e-16	***
hour21	114.00143	7.44391	15.315	< 2e-16	***
hour22	75.29220	7.45039	10.106	< 2e-16	***
hour23	37.22724	7.37198	5.050	4.51e-07	***
wkday2	-11.01087	4.16184	-2.646	0.008168	**
wkday3	-7.87850	4.10615	-1.919	0.055054	.
wkday4	-4.32022	4.10701	-1.052	0.292868	
wkday5	-2.93373	4.07082	-0.721	0.471130	
wkday6	NA	NA	NA	NA	
wkday7	16.38078	3.99532	4.100	4.17e-05	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 100.5 on 8659 degrees of freedom

Multiple R-squared: 0.6936, Adjusted R-squared: 0.6919

F-statistic: 408.3 on 48 and 8659 DF, p-value: < 2.2e-16

Start: AIC=80333.49

count ~ season + holiday + workingday + weather + temp + humidity +
windspeed + year + month + hour + wkday

Step: AIC=80333.49

count ~ season + holiday + weather + temp + humidity + windspeed +
year + month + hour + wkday

Step: AIC=80333.49

count ~ holiday + weather + temp + humidity + windspeed + year +
month + hour + wkday

	Df	Sum of Sq	RSS	AIC
- holiday	1	1816	87400849	80332
<none>			87399033	80333
- windspeed	1	143031	87542064	80346
- wkday	6	258475	87657508	80347
- humidity	1	954493	88353526	80426
- temp	1	2295989	89695022	80557
- weather	3	2571101	89970134	80580
- month	11	5734773	93133806	80865
- year	1	15873198	103272230	81785
- hour	23	100736154	188135187	86964

Step: AIC=80331.67

```
count ~ weather + temp + humidity + windspeed + year + month +  
      hour + wkday
```

	Df	Sum of Sq	RSS	AIC
<none>			87400849	80332
- windspeed	1	142973	87543822	80344
- wkday	6	265216	87666065	80346
- humidity	1	955401	88356250	80424
- temp	1	2294225	89695074	80555
- weather	3	2569967	89970816	80578
- month	11	5741240	93142089	80864
- year	1	15873066	103273914	81783
- hour	23	100753472	188154320	86963

Call:

```
lm(formula = count ~ weather + temp + humidity + windspeed +  
    year + month + hour + wkday, data = train_subset)
```

Coefficients:

(Intercept)	weather2	weather3	weather4	temp	humidity
-84.6740	-12.5972	-71.1021	-174.2638	5.0256	-0.7636
windspeed	year2012	month2	month3	month4	month5
-0.5423	86.9553	11.3517	29.6681	50.7423	80.3289
month6	month7	month8	month9	month10	month11
67.7000	40.1953	49.8426	76.9729	96.7900	76.7454
month12	hour1	hour2	hour3	hour4	hour5
75.7494	-11.4051	-24.0201	-37.4611	-38.0175	-23.4766
hour6	hour7	hour8	hour9	hour10	hour11
36.5740	170.5120	311.3926	164.7376	113.5519	140.8194
hour12	hour13	hour14	hour15	hour16	hour17
177.9067	177.2196	162.0334	168.3269	231.4745	387.3397
hour18	hour19	hour20	hour21	hour22	hour23
360.4454	245.6220	164.0973	114.0212	75.2779	37.2276
wkday2	wkday3	wkday4	wkday5	wkday6	wkday7
3.2566	6.8380	10.3442	11.7749	14.6134	16.3761

Call:

```
lm(formula = count ~ ., data = train_subset)
```

Residuals:

Min	1Q	Median	3Q	Max
-351.68	-61.47	-7.12	50.93	438.13

Coefficients: (4 not defined because of singularities)

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-84.59331	9.33850	-9.059	< 2e-16 ***
season2	67.39603	7.87769	8.555	< 2e-16 ***
season3	76.82918	7.64092	10.055	< 2e-16 ***
season4	75.49214	5.62907	13.411	< 2e-16 ***
holiday1	11.78025	7.81568	1.507	0.131781
workingday1	14.71109	4.06430	3.620	0.000297 ***
weather2	-12.58747	2.67528	-4.705	2.58e-06 ***
weather3	-71.11945	4.47545	-15.891	< 2e-16 ***
weather4	-174.84435	100.78681	-1.735	0.082813 .
temp	5.03023	0.33352	15.082	< 2e-16 ***
humidity	-0.76332	0.07849	-9.724	< 2e-16 ***
windspeed	-0.54238	0.14408	-3.764	0.000168 ***
year2012	86.95570	2.19273	39.656	< 2e-16 ***
month2	11.13212	5.39034	2.065	0.038934 *
month3	29.41766	5.77185	5.097	3.53e-07 ***
month4	-16.76422	5.99977	-2.794	0.005215 **
month5	12.64662	5.45402	2.319	0.020431 *
month6	NA	NA	NA	NA
month7	-36.80569	5.57965	-6.596	4.46e-11 ***
month8	-27.30585	5.45357	-5.007	5.64e-07 ***
month9	NA	NA	NA	NA
month10	21.16010	5.77289	3.665	0.000248 ***
month11	1.14471	5.35424	0.214	0.830712
month12	NA	NA	NA	NA
hour1	-11.39400	7.48756	-1.522	0.128115
hour2	-24.02145	7.45155	-3.224	0.001270 **
hour3	-37.44770	7.55542	-4.956	7.32e-07 ***
hour4	-38.01239	7.44329	-5.107	3.34e-07 ***
hour5	-23.47057	7.47555	-3.140	0.001697 **
hour6	36.59158	7.41431	4.935	8.15e-07 ***
hour7	170.52864	7.39633	23.056	< 2e-16 ***
hour8	311.38508	7.44159	41.844	< 2e-16 ***
hour9	164.73930	7.38202	22.316	< 2e-16 ***
hour10	113.53297	7.46630	15.206	< 2e-16 ***
hour11	140.80547	7.50670	18.757	< 2e-16 ***
hour12	177.90103	7.56740	23.509	< 2e-16 ***
hour13	177.19756	7.66452	23.119	< 2e-16 ***
hour14	162.02489	7.65093	21.177	< 2e-16 ***
hour15	168.32943	7.59391	22.166	< 2e-16 ***
hour16	231.47403	7.62564	30.355	< 2e-16 ***
hour17	387.32373	7.66555	50.528	< 2e-16 ***
hour18	360.43568	7.58264	47.534	< 2e-16 ***
hour19	245.60360	7.43058	33.053	< 2e-16 ***
hour20	164.11798	7.51229	21.847	< 2e-16 ***
hour21	114.00143	7.44391	15.315	< 2e-16 ***
hour22	75.29220	7.45039	10.106	< 2e-16 ***
hour23	37.22724	7.37198	5.050	4.51e-07 ***

wkday2	-11.01087	4.16184	-2.646	0.008168	**
wkday3	-7.87850	4.10615	-1.919	0.055054	.
wkday4	-4.32022	4.10701	-1.052	0.292868	
wkday5	-2.93373	4.07082	-0.721	0.471130	
wkday6	NA	NA	NA	NA	
wkday7	16.38078	3.99532	4.100	4.17e-05	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 100.5 on 8659 degrees of freedom

Multiple R-squared: 0.6936, Adjusted R-squared: 0.6919

F-statistic: 408.3 on 48 and 8659 DF, p-value: < 2.2e-16


Warning message in predict.lm(lm_fit, train_subset):

prediction from a rank-deficient fit may be misleadingWarning message in

predict.lm(lm_fit, test_subset):


prediction from a rank-deficient fit may be misleadingWarning message in

predict.lm(lm_fit, bike_test):





```
[24]: plot(lm_fit)
```


Warning message:
not plotting observations with leverage one:



Warning message:
not plotting observations with leverage one:







```
[12]: mse(y_act_train,y_pred_train)
      rmse(y_act_train,y_pred_train)
      rmsle(y_act_train,y_pred_train)

      mse(y_act_test,y_pred_test)
      rmse(y_act_test,y_pred_test)
      rmsle(y_act_test,y_pred_test)
```

```
9818.80878274557
99.0899025266731
1.02779704287991
9763.82893890566
98.8120890321911
1.00488949037201
```

```

[13]: # 5b. Random Forest
      Ntree=500
      Mtry = 5
      myImportance = TRUE

      # Predict Casual Counts
      set.seed(1)
      CasualData <- subset(train, select = -c(count, registered, date, atemp))
      CasualFit <- randomForest(casual ~ ., data=CasualData, ntree=Ntree,
      ↪mtry=Mtry,
                                importance=myImportance)

      # Predict Registered Counts
      RegisteredData <- subset(train, select = -c(count, casual, date, atemp))
      RegisteredFit <- randomForest(registered ~ ., data=RegisteredData,
      ↪ntree=Ntree, mtry=Mtry,
                                importance=myImportance)

[21]: varImpPlot(CasualFit)
      varImp(CasualFit)


      varImpPlot(RegisteredFit)
      varImp(RegisteredFit)

      #Inference - Casual Fit: season, holiday, windspeed and weather are not
      ↪much significant here.
      #Inference - Registered Fit: season, holiday, windspeed and weekday are not
      ↪much significant here.

```


	Overall
season	17.24307
holiday	20.74985
workingday	37.09054
weather	26.53064
temp	67.45059
humidity	63.02898
windspeed	25.46503
year	101.47007
month	39.38764
hour	192.87762
wkday	28.16724

CasualFit



	Overall
season	31.02606
holiday	17.26589
workingday	58.48382
weather	52.71161
temp	45.31425
humidity	47.63264
windspeed	19.45475
year	253.04092
month	60.28591
hour	323.42262
wkday	39.65462

RegisteredFit



```
[23]: casualFitFinal <- randomForest(casual ~ hour + year + humidity + month + temp +
  ↳workingday + wkday,
                                     data=CasualData, ntree=Ntree,
  ↳mtry=Mtry,importance=myImportance)
  RegisteredFitFinal <- randomForest(registered ~ hour + year + month +
  ↳weather + workingday + humidity + temp,
                                     data=RegisteredData, ntree=Ntree,
  ↳mtry=Mtry,importance=myImportance)
```

```
[25]: # Prediction on train data

      # Prediction on train data - casual users
      PredTrainCasual = round(predict(CasualFit, train),0)
      PredTrainCasualFinal = round(predict(casualFitFinal, train),0)
```

```

# Prediction on train data - Registered users
PredTrainRegistered = round(predict(RegisteredFit, train),0)
PredTrainRegisteredFinal = round(predict(RegisteredFitFinal,
→train),0)

# Sum up Casual and Registered to get Total Count
PredTrainCount = PredTrainCasual+PredTrainRegistered
PredTrainCountFinal = PredTrainCasualFinal+PredTrainRegisteredFinal

# Calculate Train RMSLE
rf_train_rmsle_full = rmsle(train$count, PredTrainCount)
rf_train_rmsle2_reduced = rmsle(train$count, PredTrainCountFinal)

# Prediction on test data
# Prediction on test data - casual users
PredTestCasual = round(predict(CasualFit, test),0)
PredTestCasualFinal = round(predict(casualFitFinal, test),0)

# Prediction on test data - registered users
PredTestRegistered = round(predict(RegisteredFit, test),0)
PredTestRegisteredFinal = round(predict(RegisteredFitFinal, test),0)

# Sum up Casual and Registered to get Total Count
PredTestCount = PredTestCasual+PredTestRegistered
PredTestCountFinal = PredTestCasualFinal+PredTestRegisteredFinal

# Calculate Train RMSLE
rf_test_rmsle_full = rmsle(test$count, PredTestCount)
rf_test_rmsle2_reduced = rmsle(test$count, PredTestCountFinal)

```

```

[26]: cat("Training RMSLE - Linear Regression: ", lm_train_RMSLE)
cat("\nTraining RMSLE - Random Forest (Full Model): ", rf_train_rmsle_full)
cat("\nTraining RMSLE - Random Forest (Reduced Model): : ",
→rf_train_rmsle2_reduced)

cat("\n\nTest RMSLE - Linear Regression: ", lm_test_RMSLE)
cat("\nTest RMSLE - Random Forest (Full Model): ", rf_test_rmsle_full)
cat("\nTest RMSLE - Random Forest (Reduced Model): ", rf_test_rmsle2_reduced)

```

```

Training RMSLE - Linear Regression: 1.027797
Training RMSLE - Random Forest (Full Model): 0.2561525
Training RMSLE - Random Forest (Reduced Model): : 0.2147757

```


```

Test RMSLE - Linear Regression: 1.004889
Test RMSLE - Random Forest (Full Model): 0.4212346

```


```
[27]: hist(bike$count, main="Training Data")
      hist(lm_results, main="Linear Regression Fit")
      hist(rf_results, main="Random Forest Fit")

      # Inference: The distribution of predicted count looks similar to that
      ↳ of train data.
```



```
Error in hist(rf_results, main = "Random Forest Fit"): object
↳ 'rf_results' not found
Traceback:
```

```
1. hist(rf_results, main = "Random Forest Fit")
```



```
[49]: # Save the RF results
      rf_test_casual = round(predict(casualFitFinal, bike_test),0)
      rf_test_registered = round(predict(RegisteredFitFinal, bike_test),)
      rf_results = rf_test_casual + rf_test_registered
```

```
[36]: gbmtree=4000
      iDepth = 3
      set.seed(1)

      # Predict Casual Counts
      CasualData <- subset(train, select = -c(count, registered, atemp, date))
```



```
gbm.Casual <- gbm(log1p(casual)~.,data=CasualData,distribution= "gaussian",n.
→trees=gbmtree,interaction.depth=iDepth)
```

```
# Predict Registered Counts
```


```
RegisteredData <- subset(train, select = -c(count, casual, atemp, date))
gbm.Registered <- gbm(log1p(registered)~.,data=RegisteredData,distribution=
→"gaussian",n.trees=gbmtree,interaction.depth=iDepth)
```

```
[37]: summary(gbm.Casual)
summary(gbm.Registered)
```


```
##Inference - gbm Casual: season, holiday, year, windspeed are not much
→significant here.
```

```
##Inference - gbm Registered: holiday, windspeed, season, weather are not much
→significant here.
```

	var	rel.inf
hour	hour	57.85237775
temp	temp	16.11053055
month	month	10.05511534
humidity	humidity	4.15202731
wkday	wkday	4.05126767
workingday	workingday	3.52722769
weather	weather	1.90563745
windspeed	windspeed	1.10793870
year	year	1.02249663
holiday	holiday	0.19081563
season	season	0.02456529



	var	rel.inf
hour	hour	73.9936471
month	month	5.7223401
workingday	workingday	5.0367884
year	year	3.4921877
wkday	wkday	3.4287371
temp	temp	2.7665790
humidity	humidity	2.5423855
weather	weather	1.4728776
season	season	0.8166783
windspeed	windspeed	0.5150276
holiday	holiday	0.2127515



```
[38]: gbm.CasualFinal <- gbm(log1p(casual) ~ hour + workingday + temp + month +
  ↳ wkday + humidity + weather,
                                     data=CasualData, distribution= "gaussian",n.
  ↳ trees=gbmtree,interaction.depth=iDepth)
gbm.RegisteredFinal <- gbm(log1p(registered) ~ hour + year + workingday + month
  ↳ + wkday + humidity + temp,
                                     data=RegisteredData, distribution= "gaussian",n.
  ↳ trees=gbmtree,interaction.depth=iDepth)
```

```
[39]: # Prediction on train data
      # Prediction on train data - casual users

gbm.PredTrainCasual <- predict(gbm.Casual, train, n.trees=gbmtree)
gbm.PredTrainCasualFinal <- predict(gbm.CasualFinal, train, n.trees=gbmtree)
```

```

# Prediction on train data - Registered users
gbm.PredTrainRegistered <- predict(gbm.Registered, train, n.trees=gbmtree)
gbm.PredTrainRegisteredFinal <- predict(gbm.RegisteredFinal, train, n.
  →trees=gbmtree)

# Sum up Casual and Registered to get Total Count
gbm.PredTrainCount <- round(exp(gbm.PredTrainCasual) - 1, 0) + round(exp(gbm.
  →PredTrainRegistered) - 1, 0)
gbm.PredTrainCountFinal <- round(exp(gbm.PredTrainCasualFinal) - 1, 0) +
  →round(exp(gbm.PredTrainRegisteredFinal) - 1, 0)

# Calculate Train RMSLE
gbm.rf_train_rmsle_full <- rmsle(train$count, gbm.PredTrainCount)
gbm.rf_train_rmsle2_reduced <- rmsle(train$count, gbm.PredTrainCountFinal)

# Prediction on test data
# Prediction on test data - casual users
gbm.PredTestCasual = predict(gbm.Casual, test, n.trees=gbmtree)
gbm.PredTestCasualFinal = predict(gbm.CasualFinal, test, n.trees=gbmtree)

# Prediction on test data - registered users
gbm.PredTestRegistered = predict(gbm.Registered, test, n.trees=gbmtree)
gbm.PredTestRegisteredFinal = predict(gbm.RegisteredFinal, test, n.
  →trees=gbmtree)

# Sum up Casual and Registered to get Total Count
gbm.PredTestCount = round(exp(gbm.PredTestCasual) - 1, 0) + round(exp(gbm.
  →PredTestRegistered) - 1, 0)
gbm.PredTestCountFinal = round(exp(gbm.PredTestCasualFinal) - 1, 0) +
  →round(exp(gbm.PredTestRegisteredFinal) - 1, 0)

# Calculate Test RMSLE
gbm.rf_test_rmsle_full = rmsle(test$count, gbm.PredTestCount)
gbm.rf_test_rmsle2_reduced = rmsle(test$count, gbm.PredTestCountFinal)

```

```

[75]: gbm.rf_train_rmsle_full
      gbm.rf_train_rmsle2_reduced

gbm.rf_test_rmsle_full
gbm.rf_test_rmsle2_reduced

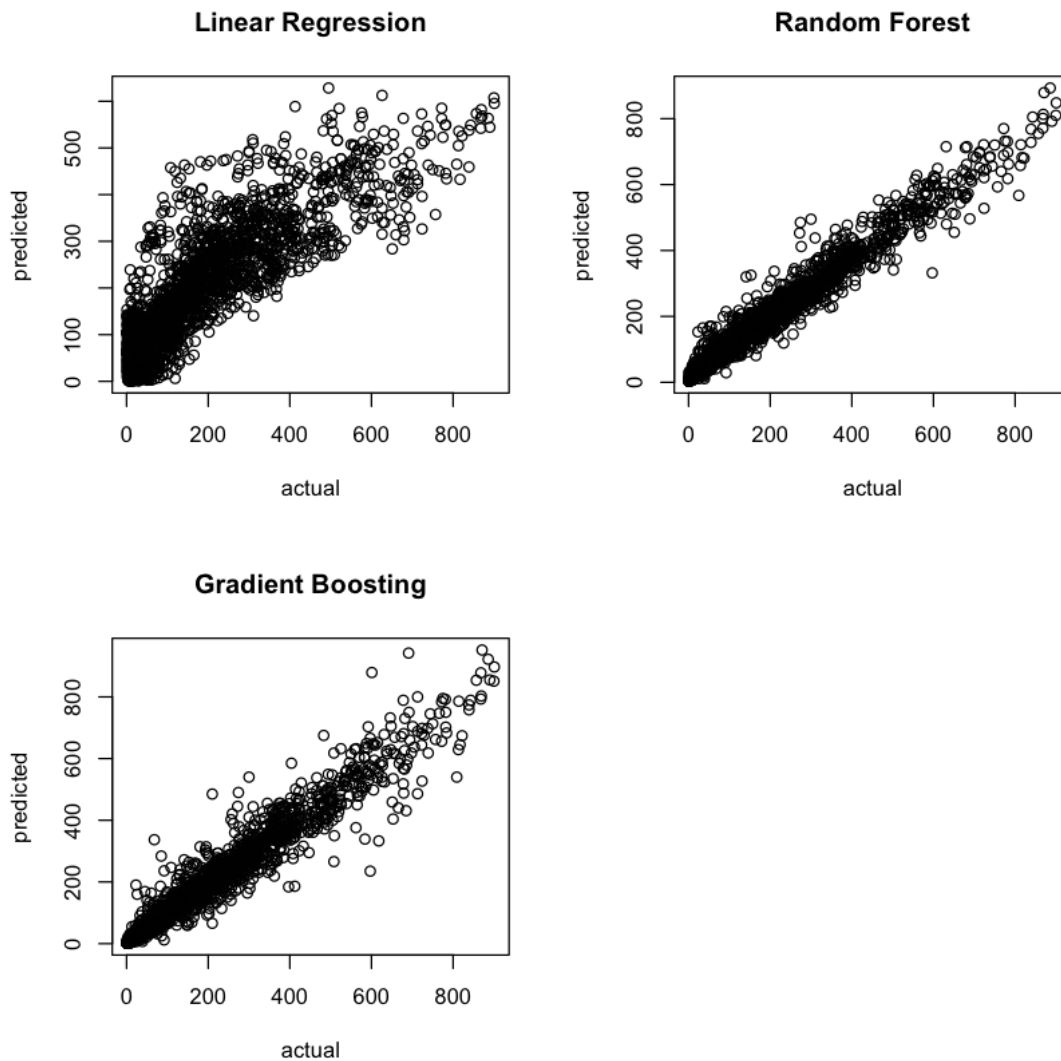
```

```

0.215991634854015
0.241986515661299
0.277804015565458
0.298095056429441

```

```
[76]: par(mfrow=c(2,2))
plot(y_act_test, y_pred_test, main="Linear Regression", xlab="actual",
      ylab="predicted")
plot(test$count, PredTestCount, main="Random Forest", xlab="actual",
      ylab="predicted")
plot(test$count, gbm.PredTestCountFinal, main="Gradient Boosting",
      xlab="actual", ylab="predicted")
```



```
[78]: cor(y_act_test, y_pred_test)
cor(test$count, PredTestCountFinal)
cor(test$count, gbm.PredTestCountFinal)
```

```
0.846726206103257
0.975317861770545
```

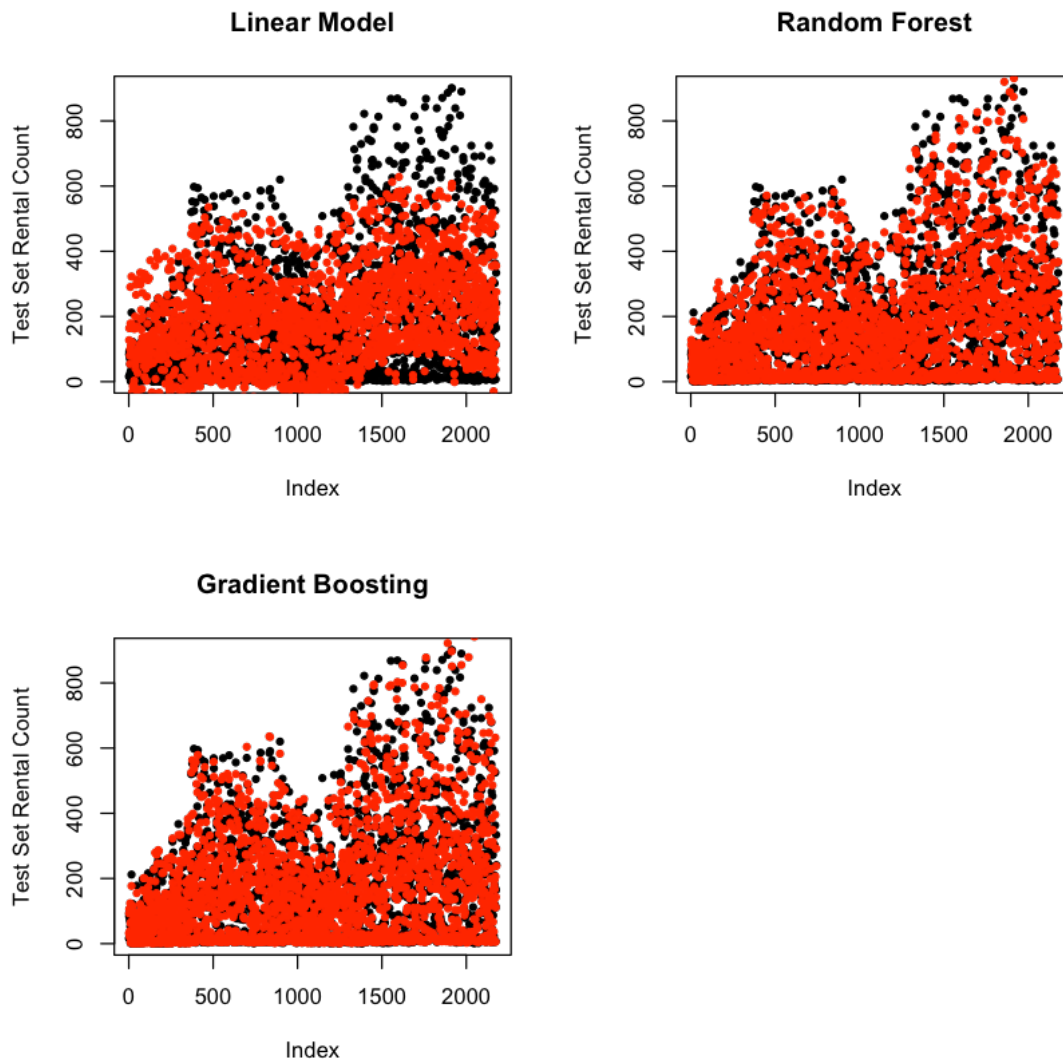
0.970021541913715

```
[79]: par(mfrow=c(2,2))
plot(test_subset$count, main = "Linear Model", ylab = "Test Set Rental Count",
      pch = 20)
points(predict(lm_fit, newdata = test), col = "red", pch = 20)

plot(test_subset$count, main = "Random Forest", ylab = "Test Set Rental Count",
      pch = 20)
points(PredTestCountFinal, col = "red", pch = 20)





plot(test_subset$count, main = "Gradient Boosting", ylab = "Test Set Rental
      Count", pch = 20)
points(gbm.PredTestCountFinal, col = "red", pch = 20)
```

Warning message in predict.lm(lm_fit, newdata = test):



```
[47]: # Save the RF results
      gbm_test_casual = round(predict(gbm.CasualFinal, bike_test, n.
      ↪trees=gbmtree),0)
      gbm_test_registered = round(predict(gbm.RegisteredFinal, bike_test, n.
      ↪trees=gbmtree),0)
      gbm_results = gbm_test_casual + gbm_test_registered
```

```
[51]: par(mfrow=c(2,2))
hist(bike$count, main="Training Data")
hist(lm_results, main="Linear Regression Fit")
hist(rf_results, main="Random Forest Fit")
hist(rf_results, main="Gradient Boosting Fit")
```



[: